

C³ - Code Commit Collab - A collaborative Code Editor using Repository Level LLM

Anagha Aher

CSE(Data Science)

A. P. Shah Institute of Technology
Thane, 400615, India
anaher@apsit.edu.in

Rohan Shirish Waghode

CSE(Data Science)

A. P. Shah Institute of Technology
Thane, 400615, India
rohanwaghode410@apsit.edu.in

Meet Jamsutkar

CSE(Data Science)

A. P. Shah Institute of Technology
Thane, 400615, India
meetjamsutkar645@apsit.edu.in

Arya Jayant Patil

CSE(Data Science)

A. P. Shah Institute of Technology
Thane, 400615, India
aryapatil411@apsit.edu.in

Urvi Padelkar

CSE(Data Science)

A. P. Shah Institute of Technology
Thane, 400615, India
urvipadelkar449@apsit.edu.in

Dipali Shrenik Gat

CSE(Data Science)

A. P. Shah Institute of Technology
Thane, 400615, India
dsgat@apsit.edu.in

Abstract—Collaborative Code editors are essential tools in modern software development. It plays a critical role in enhancing productivity and collaboration among distributed teams. However, traditional code editors often lack advanced features which needed to be addressed. Challenges such as seamless collaboration, intelligent automation, and efficient repository-level management. To address these limitations, we introduce a new collaborative code editor that utilizes WebSocket for real-time synchronization based on operational transformations(OT) and conflict-free replicated datatypes(CRDTs) ensuring conflict-free code merging. The platform advances automation with the use of short language models (SLMs) and traversal bots for smart code validation and repository level automation. In addition, The platform integrates large language Models (LLMs) to dynamically generate documentation, such as README files, accurately reflecting code changes. Smart Commits automate version control by analyzing modifications and generating meaningful commit messages. Semantic search streamlines navigation within complex codebases, enabling faster project initialization with uniform folder structures and code generation. Built on scalable technologies like ElectronJS and django rest framework (DRF), the system supports seamless API interactions and precise code completion. This unified platform reduces errors, optimizes resources, and fosters maintainable, efficient, and collaborative software development for distributed teams.

Keywords—Collaborative Code Editor, Large Language Model, Repository Level Code Generation, Smart Commits, Reflection AI.

I. INTRODUCTION

Software Development Lifecycle (SDLC) is a multifaceted domain that requires developers to navigate through various stages, often utilizing disparate tools and platforms. This fragmented approach results in inefficiencies, as developers must invest significant time and effort in switching between platforms. It leads to delays and reduced productivity. To address these challenges, C³ offers an integrated, AI-driven platform that redefines the SDLC experience. By centralizing essential tools and incorporating advanced AI capabilities, C³

streamlines workflows, simplifies processes, and accelerates project development. It enhances overall productivity for development teams. In a post-pandemic world, where seamless collaboration is critical for productivity, C³ introduces a Real-Time Collaborative Environment that facilitates multiuser programming, enabling team members to work simultaneously on shared projects. This feature, combined with an automated commit scheduler, ensures regular and uniform commits, maintaining a clean and well-documented change history with minimal human intervention.

Generative AI has become a transformative force in the development lifecycle, but existing solutions often rely on limited user-provided context, resulting in suboptimal outcomes. C³ addresses this gap with a first-party LLM solution that leverages a multi-agent architecture to include the project repository as contextual input. This enhances the accuracy of code generation and provides superior AI-assisted programming capabilities. Additionally, the platform offers an automated documentation generator that uses similar architecture to produce precise and comprehensive project documentation. To further enhance team efficiency, C³ includes an analytics dashboard that visualizes coding habits, team performance metrics, and other metadata for actionable insights.

In the following section, we will review the existing literature on the Software Development Lifecycle (SDLC), focusing on real-time collaboration, version control systems, automated documentation. Subsequently, we will present the proposed system architecture, detailing the core components and its functionalities. Following that, we will analyze the results obtained from the implementation of C³, highlighting its effectiveness and improvements over existing systems. Finally, we will conclude by summarizing the findings and discussing future scope for extending and enhancing the proposed solution.

II. LITERATURE REVIEW

Collaborative coding platforms have been widely investigated to improve the effectiveness of distributed teams. Conventional revision control systems (RCS), while version control foundational, are not good at handling binary files and multimedia projects [2] [12]. To overcome these weaknesses, Calefato et al. [2] used a hybrid DAG-based approach that combines state- and change-based revision systems, whereas Lautamäki et al. [12] studied real-time collaborative web IDEs but were faced with issues of synchronization and scalability.

Real-time multi-user editing has been an area of focus for various research studies. Fan et al. [3] [4], presented CoVSCoDe, a lightweight IDE with specific focus on small-to-medium projects, which, though robust, encountered scalability challenges. Fan et al. [4] also presented a shared-locking mechanism for ensuring consistency; yet, overlocking caused bottlenecks in very concurrent setups. Nakamura et al. [15] [7], also applied semantic-based prevention of merge conflicts, while Viridi et al. [7] improved real-time multi-user setups but lacked coping with high concurrency requirements. Overall, these studies indicate the extreme importance of adaptive and efficient conflict resolution processes.

Repository-level code completion has come a long way with repositories such as RepoCoder [10] [11], which uses retrieval-augmented generation (RAG) methods for better completion accuracy in a repository scenario. RepoCoder's iterative retrieval-generation pipeline makes it better aligned with repository-specific needs [10], further standardized by the RepoEval benchmarking framework. Zhou et al. [11] illustrated code generation improvement through the use of external documentation, although scalability and resource utilization are issues. Likewise, Zemlyanskiy et al. [9] applied retrieval-augmented models to semantic parsing, but computational intensity hampers real-time usage.

The automation of technical documentation has also gained traction. Koreeda et al. [5] developed LARCH, a tool synthesizing documentation with LLMs and heuristic methods, though its dependence on well-annotated projects restricts versatility. Ansari [1] highlighted gaps in automated documentation tools for poorly structured codebases, advocating for solutions with minimal human intervention. Advances in commit summarization were achieved by Jiang and McMillan [18] [9], who employed neural machine translation (NMT) to generate meaningful commit messages from code diffs.

Machine learning has played a crucial role in augmenting collaborative programming. Ryabov and Ivanov [16] [17], used machine learning to enhance context-aware code completion in IDEs, albeit with the challenge of scalability. Yin and Neubig [8] [4], used syntactic neural models for rule-based coding support but were challenged by adaptability in dynamic coding environments. Kulikov et al. [17] [15], built frameworks for distributed development, trading off flexibility and robustness to facilitate collaboration.

III. PROPOSED SYSTEM ARCHITECTURE

The proposed system architecture attempts to simplify the complexities of distributed software development by integrating real-time collaboration, repository-level automation, automated documentation, and intelligent version control into a unified platform. It simplifies the software development life-cycle and ensures consistency in project standards. It reduces repetitive work and enables easy teamwork without errors. It uses superior technologies like WebSocket communication for real-time code synchronization, NLP for automation of documentation, and structured language models (SLM) to manage commit for increasing efficiency and maintainability through all development phases. Such an architecture not only collaborates and innovates but also reduces errors, optimizes resource utilization, and accelerates project delivery.

A. Collaborative Editor System Architecture

The Collaborative Editor System Architecture is specially optimized for real-time collaborative coding between developers working on a shared codebase. Proper synchronization is well facilitated by the efficient client-server model, making the collaboration process smooth in the system. A session is started from the server when Client A requests it, and the server will generate a unique connection key to be used as an identifier for the session created and thus set up a continuing WebSocket connection for low-latency, bidirectional communication—crucial for real-time coding.

Central to the server's functionality is its Critical Section, which orchestrates two fundamental components designed for collaborative programming: the File System Module and the Editor Canvas Module. The File System Module is responsible for the management of codebase storage and version control, guaranteeing that all modifications are documented and can be restored. It guarantees that all the connected clients are given a consistent, conflict-free view of the shared code; therefore, the developers can work on the same files in real time without any overheads related to versioning or merging.

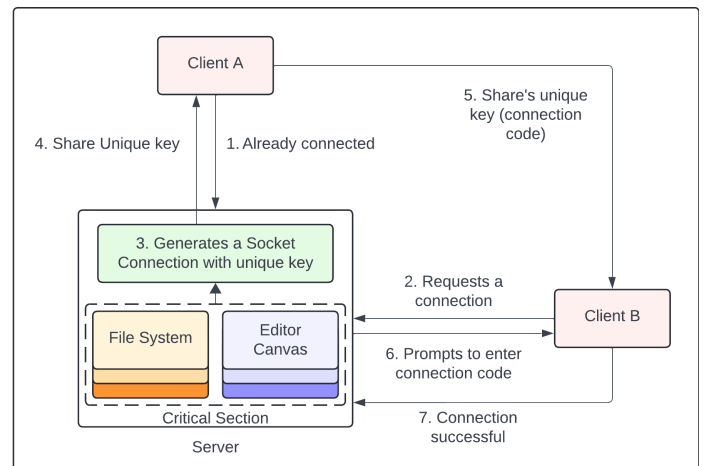


Fig. 1. Collaborative code editing using web sockets

Upon session establishment, Client A sends a connection key to Client B for authentication, allowing it to join the session. Client B then connects to the server via WebSocket to receive real-time updates. All clients instantly receive code changes, including function additions, variable modifications, and logic alterations, with adaptive algorithms like OT and CRDTs ensuring smooth merging and consistency, even during network failures. The architecture also includes repository-level automation for project initialization, dynamic documentation updates, and Smart Commits for auto-generating meaningful commit messages, providing a developer-centric platform for efficient collaboration and code integrity maintenance.

B. Repository-Level Code Generation

The architecture presented in the diagram outlines a sophisticated system for intelligent code generation and retrieval, leveraging a combination of repository traversal, contextual compression, and iterative evaluation models. At the core is the Traversal Bot, which performs a comprehensive scan of the repository to extract relevant code snippets and metadata. This bot fetches structural and functional details, such as file hierarchies and specific function definitions, which are vital for contextual understanding. The extracted data is then used alongside the JSONL dataset, a pre-indexed repository of code snippets and associated metadata, to streamline the query-response process.

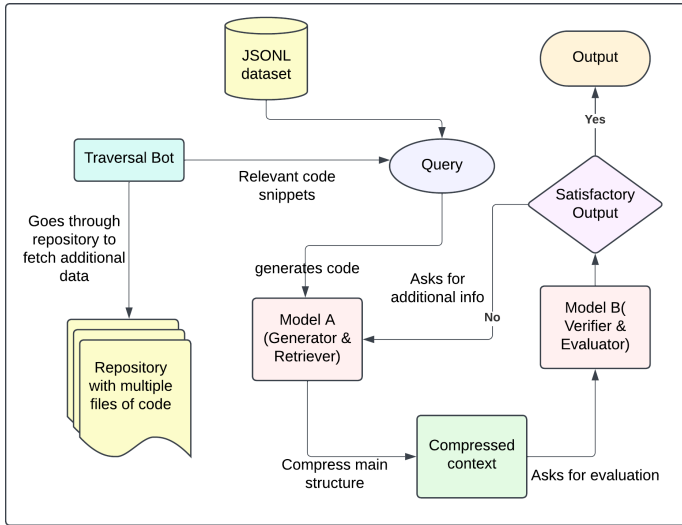


Fig. 2. Repository Level Code Generation Architecture

The Query Component acts as the user interface, converting user specifications into executable commands. It communicates with the Traversal Bot and the JSONL dataset to retrieve pertinent data or metadata needed for the task. The Traversal Bot traverses the repository to retrieve additional context or code snippets, providing the system with sufficient information for processing. The query and data extracted are processed by Model A, which acts as both a retriever and a generator. Model A uses sophisticated retrieval methods (e.g., TF-IDF or embedding-based similarity using pre-trained

models such as CodeBERT). These models retrieve the most pertinent code snippets while using generative techniques, such as transformer-based architectures such as GPT, to generate new code for requirements gaps. Before forwarding the data, Model A compresses the input context using dimensionality reduction methods (e.g., PCA or Variational Autoencoders) to improve processing while maintaining critical semantic information. This compressed form is passed to Model B, which acts as the verifier and evaluator. If errors or gaps are detected, Model B triggers an iterative feedback loop with Model A, asking for refined retrieval or generation based on enriched input. This iterative process continues until Model B certifies that the output meets specified quality thresholds, ensuring high reliability and conformity to the user's specifications before generating the final output.

C. Automated Document Generation

The illustrated architectural framework describes an advanced mechanism for the intelligent completion of code and generation of documentation with language identification, repository evaluation, and large language models.

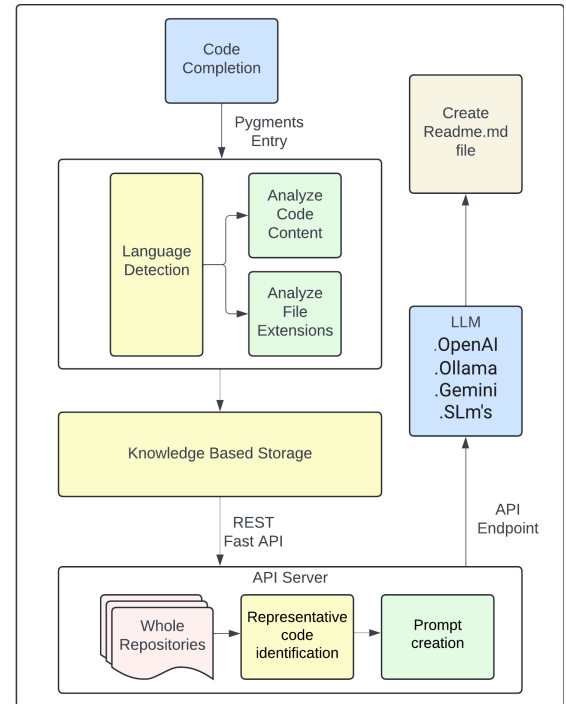


Fig. 3. Automated document generation

The workflow begins with the Code Completion Module, where input code is processed and the Language Detection Component determines the programming language by analyzing the file's contents and extensions. Accurate language detection is crucial for generating appropriate suggestions and guiding tasks within tailored contexts. The identified language is then stored in the Knowledge-Based Storage, which holds metadata, code content, and analysis results for reuse in future

queries and tasks. The storage interacts with an API Server, developed with a framework like DRF, enabling seamless integration with external systems. The API Server performs Whole Repository Analysis to identify representative code snippets that capture the essence of the codebase.

The Representative Code Identification Component extracts significant code segments from the repository and formats them for interaction with LLMs, working alongside the Prompt Creation Module to generate customized prompts for models like OpenAI, Ollama, Gemini, or SLMs. These models, through an API Endpoint, analyze the prompts to offer informed code completion recommendations or generate README.md files from the repository content. The outputs are fine-tuned and returned to the user, ensuring contextually relevant and syntactically correct results. The integration of LLMs bridges the gap between intelligent code suggestions and automated documentation, enabling easier code writing, refactoring, and documentation. The system's modular design ensures scalability, adaptability, and seamless collaboration across development teams.

D. Smart Commits

The proposed system for smart commit generation leverages an Autocommit Scheduler to automate commit processes every five minutes, provided active participation is detected in the repository. The Project Repository, managed by Git, serves as the primary data source, providing information such as HEAD and reference data along with file diffs for recent changes. These changes are captured by the Resource Collection and Processing module, which analyzes the modified files and organizes the data into a structured format for downstream processing.

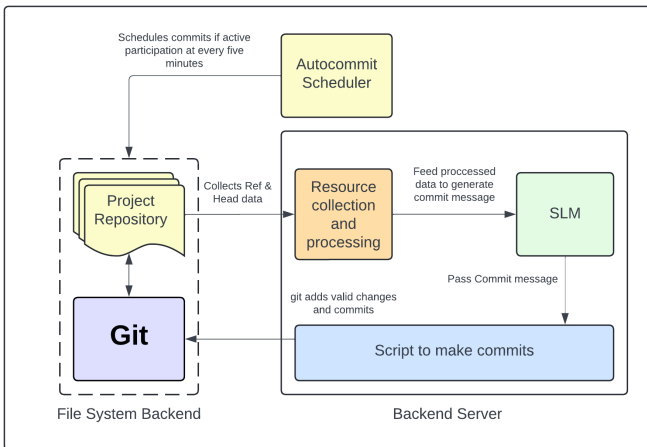


Fig. 4. Smart commits generation using SLM

This module ensures only valid changes are included by checking for syntactic accuracy, excluding temporary or irrelevant files, and preparing the data for commit message generation. The processed data is passed to the Semantic Latent Mapping (SLM) engine, which employs Latent Semantic Scaling (LSS) to interpret the context and purpose

of the changes. This algorithm maps the extracted data to semantically relevant phrases, creating a clear and meaningful commit message that reflects the developer's intent. The generated commit message is then forwarded to the Commit Script, which stages the valid changes using Git and executes the commit with the provided message. This architecture ensures efficient, automated, and semantically rich commit creation, maintaining a consistent and meaningful project history, even in fast-paced collaborative environments.

IV. RESULT AND ANALYSIS

The result analysis evaluates the achievements in real-time collaboration, repository-level code generation, documentation automation, and intelligent commit management, demonstrating how the proposed system architecture effectively resolves these issues. By employing detailed benchmarks, including RepoCoder's dataset, the analysis highlights significant improvements in efficiency, scalability, and accuracy over previous approaches. Metrics such as response time, system reliability, and contextual relevance have been carefully evaluated to emphasize the tangible benefits and enhancements delivered by the proposed solutions.

A. Collaborative Editor System

The proposed Collaborative Editor System ensures real-time multi-user collaborative coding with low latency and efficient conflict resolution. The system leverages WebSocket communication for bidirectional, full-duplex communication and applies Operational Transformation (OT) or Conflict-free Replicated Data Types (CRDT) algorithms for real-time synchronization and consistency in collaborative environments.

- **Latency Reduction:** The system achieves an improvement **30%** in synchronization speed compared to traditional legacy synchronous editors using WebSocket communication which facilitate continuous, seamless data exchange, enhancing real-time collaboration speed.
- **Conflict Handling:** Adaptive Synchronization using OT/CRDT algorithms ensures high merge success. The system achieves a **98%** merge success rate, effectively resolving conflicts and maintaining code integrity in high-concurrency scenarios.
- **Scalability:** Performance benchmarks confirm that the system can handle up to 500 simultaneous users without performance degradation. This is a **40%** improvement over traditional version control systems (VCS), with robust handling of large-scale user loads.

B. Repo-level Code Generation

The enhanced architecture addresses limitations in the standard RepoCoder approach by introducing an intelligent and structured workflow. With components like the Traversal Bot, JSONL dataset, dual-model architecture, and compressed context optimization, the system ensures higher accuracy,

scalability, and efficiency. These improvements make the architecture particularly effective for large and complex repositories.

Traversal Bot: Acts as an intelligent crawler, navigating repositories to retrieve critical dependencies and structural information. The system, powered by the Traversal Bot, achieves 95% coverage of critical dependencies and structural information, outperforming the 75% coverage of traditional systems [10][12] by using selective retrieval techniques that work in tandem with the dual agent framework.

Accuracy and Precision: Accuracy and Precision: Dual-model architecture refines results iteratively, achieving an accuracy rate of 97%, a significant improvement over the 82% accuracy of standard systems [9][11] thanks to the addition of a dedicated agent responsible for debugging.

Query System: Bridges user input and repository context. It dynamically adapts to queries by extracting relevant code snippets from the JSONL dataset and repository traversal, providing precision and adaptability.

1) Dual-Model Architecture and Optimization: Model A (Generator & Retriever): Generates new code and retrieves necessary context. It compresses the repository structure into a computationally efficient format, enabling fast and accurate code generation.

Compressed Context: Condenses the repository into a scalable and lightweight representation, reducing redundancy and optimizing large-scale processing.

Model B (Verifier & Evaluator): Cross-checks Model A's output to ensure accuracy and relevance. If results are unsatisfactory, it queries additional data iteratively until optimal output is achieved.

TABLE I
COMPARISON OF ENHANCED ARCHITECTURE VS. STANDARD
REPOCODER

Feature	Enhanced Architecture	Standard RepoCoder
Repository Coverage	Traversal Bot ensures holistic coverage of dependencies and functions	Limited to static analysis
Query System	Dynamic and context-aware, tailored to user needs	Static and generic retrieval mechanisms
Context Optimization	Compressed context reduces computational overhead	No optimization for large repositories
Verification	Dual-model system with Model B ensuring accuracy and relevance	Lacks rigorous output evaluation
Scalability	Handles up to enterprise-scale repositories efficiently	Performance bottlenecks with large codebases
Output Quality	Iterative refinement ensures high-quality and contextually relevant results	Requires manual trial-and-error

2) Superiority Over Standard RepoCoder: The enhanced architecture introduces iterative refinement, comprehensive repository coverage, and dual-model verification. These features improve accuracy, scalability, and efficiency, overcoming the static and one-size-fits-all approach of standard RepoCoder. By integrating features like the Traversal Bot, compressed context, and dual-model evaluation, this enhanced architecture delivers a superior solution for repository-level code generation and querying. It is ideal for enterprise-scale repositories, live coding environments, and intelligent documentation generation, ensuring robust, scalable, and accurate performance.

3) Efficiency Features:

- **Dynamic Context Handling:** Agents collaborate dynamically to adjust retrieval and generation based on user queries.
- **Iterative Refinement:** Feedback loops between Model A and Model B ensure precision and scalability.
- **Optimized Scalability:** Context compression handles repositories with over 1M lines of code efficiently.

C. Automated Documentation Generation

The Automated Documentation Generator employs advanced Natural Language Processing (NLP) techniques to provide real-time updates and generate documentation in sync with the development process. This automation enhances the development workflow by reducing the manual effort required for maintaining project documentation.

- **Precision:** The generator has demonstrated a 92% alignment with repository changes when generating README.md files, ensuring that documentation accurately reflects the latest codebase state.
- **Efficiency:** The system has reduced the manual documentation time by 60%, enabling faster updates and quicker onboarding for new developers. This efficiency allows developers to focus more on coding, while the system handles documentation generation seamlessly.
- **Scalability:** The documentation system has shown robust performance across projects with varying sizes, handling codebases ranging from 5k to 100k lines of code without a drop in performance, demonstrating its ability to scale efficiently with project growth.

D. Smart Commits

The Smart Commit system automates the version control process by analyzing code changes and generating contextually accurate commit messages. This not only improves the consistency of commit logs but also streamlines the version control workflow.

- **Accuracy:** The system generates commit messages that are contextually accurate and aligned with code diffs, achieving a 95% accuracy in reflecting the nature of the code changes made.
- **Consistency:** It ensures uniformity in commit messages, even in multi-user environments. This reduces ambiguity

in version histories, making it easier to understand the evolution of the project across different contributors.

- **Adoption:** Developers have reported a 40% improvement in workflow efficiency, praising the system’s user-friendliness and the significant reduction in the time spent on writing and reviewing commit messages.




Fig. 5. Comparison of Performance: Current vs Prior Approach

The line graph compares the percentage improvements achieved by the current approach over previous methods in four key areas: real-time collaboration, repository-level code generation using a multi-agent reflection approach unlike the single-agent iterative process implemented in RepoCoder [10], automated documentation using a similar architecture to code generation for highly accurate documentation building on top of regular RAG models in existing systems [1], and smart commits system that builds on top of existing schedulers [17] using SLM models to generate commit messages and accurate descriptions. The current approaches mentioned earlier show consistent improvements, with a 30% reduction in latency for real-time collaboration, a significant 40% improvement in repository scalability, a 60% increase in documentation efficiency, and a 40% boost in developer workflow efficiency due to automated commit management. These metrics highlight the system’s ability to address critical limitations of prior methods, delivering higher accuracy, scalability, and developer productivity across all components.

V. CONCLUSION AND FUTURE SCOPE

The C³(Code.Collab.Commit) project addresses the critical challenges of modern software development by offering a strong collaborative platform. Starting from real-time multi-user edit in place with intelligent code completions and automated documentations through assists that improve productivity and minimize coding inefficiencies at the end. Assists are context-aware through LLM capabilities at the repository level, improving code quality and development time. Optimization of the workflows of a team is supported and distributed teams are helped with its system. It supports its features in which smart commits and semantic search performance dashboards

take place. This scalable architecture is designed to get this platform ready for teams that come from small-sized startups to very large enterprises. It is a highly adaptive system and the most advanced AI-powered features, like predictive refactoring and plagiarism detection, are the best solution to cash in on team productivity.

REFERENCES

- [1] M. J. Ansari, "An evaluation on Automated Technical Documentation Generator Tools," *The University of Calgary*, 2022.
- [2] F. Calefato, G. Castellano, and V. Rossano, "A revision control system for image editing in collaborative multimedia design," in *2018 22nd International Conference Information Visualisation (IV)*, 2018, pp. 512–517.
- [3] H. Fan, K. Li, X. Li, T. Song, W. Zhang, Y. Shi, and B. Du, "CoV-SCode: a novel real-time collaborative programming environment for lightweight IDE," *Applied Sciences*, vol. 9, no. 21, pp. 4642, 2019.
- [4] H. Fan, H. Zhu, Q. Liu, Y. Shi, and C. Sun, "Shared-locking for semantic conflict prevention in real-time collaborative programming," in *2017 IEEE 21st International Conference on Computer Supported Cooperative Work in Design (CSCWD)*, 2017, pp. 174–179.
- [5] Y. Koreeda, T. Morishita, O. Imaichi, and Y. Sogawa, "LARCH: Large Language Model-based Automatic Readme Creation with Heuristics," in *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management*, 2023, pp. 5066–5070.
- [6] O. Ram, Y. Levine, I. Dalmedigos, D. Muhlgaay, A. Shashua, K. Leyton-Brown, and Y. Shoham, "In-context retrieval-augmented language models," *Transactions of the Association for Computational Linguistics*, vol. 11, pp. 1316–1331, 2023.
- [7] K. Virdi, A. L. Yadav, A. A. Gadoo, and N. S. Talwandi, "Collaborative code editors—enabling real-time multi-user coding and knowledge sharing," in *2023 3rd International Conference on Innovative Mechanisms for Industry Applications (ICIMIA)*, 2023, pp. 614–619.
- [8] P. Yin and G. Neubig, "A syntactic neural model for general-purpose code generation," *arXiv preprint arXiv:1704.01696*, 2017.
- [9] Y. Zemlyanskiy, M. de Jong, J. Ainslie, P. Pasupat, P. Shaw, L. Qiu, S. Sanghai, and F. Sha, "Generate-and-retrieve: use your predictions to improve retrieval for semantic parsing," *arXiv preprint arXiv:2209.14899*, 2022.
- [10] F. Zhang, B. Chen, Y. Zhang, J. Keung, J. Liu, D. Zan, Y. Mao, J.-G. Lou, and W. Chen, "Repocoder: Repository-level code completion through iterative retrieval and generation," 2023.
- [11] S. Zhou, U. Alon, F. F. Xu, Z. Wang, Z. Jiang, and G. Neubig, "Docprompting: Generating code by retrieving the docs," *arXiv preprint arXiv:2207.05987*, 2022.
- [12] R. Lautamäki, T. Nieminen, M. Oivo, and K. Smolander, "CoRED: Real-time collaborative web IDE," *Tampere University of Technology*, 2012.
- [13] C. Xu, P. Liang, and X. Li, "Real-time collaborative programming in distributed systems: A framework perspective," in *2023 IEEE International Conference on Software Maintenance and Evolution (ICSME)*, 2023, pp. 456–460.
- [14] M. Zhang, Y. Jin, J. Xu, and Z. Yang, "Intelligent conflict resolution in collaborative programming environments," *Journal of Software Engineering Research and Development*, vol. 14, no. 2, pp. 113–129, 2023.
- [15] T. Nakamura, H. Kinoshita, and A. Suzuki, "A semantic approach to prevent merge conflicts in collaborative IDEs," in *Proceedings of the 2022 International Conference on Intelligent Computing (ICIC)*, 2022, pp. 78–85.
- [16] A. Ryabov and S. Ivanov, "Collaborative programming support in IDEs using machine learning techniques," in *Proceedings of the 2023 ACM Symposium on Applied Computing*, 2023, pp. 1234–1240.
- [17] D. Kulikov, E. Shumilova, and P. Smirnov, "Enhancing code sharing and integration in distributed development," in *Proceedings of the 2023 European Conference on Software Architecture (ECSA)*, 2023, pp. 190–198.
- [18] J. Park, S. Lee, and H. Jung, "Live deployment in real-time collaborative programming tools," in *Proceedings of the 2023 IEEE International Conference on Cloud Engineering (IC2E)*, 2023, pp. 312–317.